

NAME: M. Siri Chandana

HT.NO:2203A51492

BATCH:12

AIML ASSIGNMENT-6

Question 1:

<https://www.kaggle.com/datasets/iabhishekofficial/mobile-price-classification?select=train.csv>

From the above data:

- a) Read the data with pandas and find features and target variables
- b) Normalize the data with min-max scaling
- c) Split the data into train and test.

Question 2:

Bob has started his own mobile company. He wants to give tough fight to big companies like

Apple, Samsung etc.

He does not know how to estimate price of mobiles his company creates. In this competitive

mobile phone market you cannot simply assume things. To solve this problem he collects sales

data of mobile phones of various companies.

Bob wants to find out some relation between features of a mobile phone (eg:- RAM, Internal

Memory etc) and its selling price. But he is not so good at Machine Learning.

So he needs your

help to solve this problem. And provide the results like accuracy, precision recall and confusion

matrix.

<https://www.kaggle.com/datasets/iabhishekofficial/mobile-price-classification?select=train.csv>.

```
yp=model.predict(x_test_scaled)
```

```
from sklearn.metrics import accuracy_score, precision_score, recall_score, confusion_matrix
```

```
accuracy=accuracy_score(y_test,yp)  
print(accuracy)
```

0.95

```
precision=precision_score(y_test,yp,average='weighted')  
print(precision)
```

0.9499290394431786

```
recall=recall_score(y_test,yp,average='weighted')  
print(recall)
```

0.95

1350 3
Name: price_range, Length: 1400, dtype: int64

```
from sklearn.preprocessing import StandardScaler
```

1

```
scaler = StandardScaler()  
x_train_scaled = scaler.fit_transform(x_train)  
x_test_scaled = scaler.transform(x_test)
```

1

```
from sklearn.linear_model import LogisticRegression
```

1

```
model=LogisticRegression()  
model.fit(x_train_scaled,y_train)
```

1

```
yp=model.predict(x_test_scaled)
```

```
print(y_test)
```

```
[26]
```

```
... 423      0
     1495    0
     1618    2
     1099    0
     1307    0
     ..
     14      0
     282    1
     952    3
     1079    2
     486    2
Name: price_range, Length: 600, dtype: int64
```

```
print(y_train)
```

```
[27]
```

```
... 993      0
     1156    3
     615    2
     703    3
     1130    3
     ..
     1016    0
     165    3
     7      0
     219    3
     1350    3
```

```
print(x_test)
```

	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_memory	\
423	1681	1	2.5	0	2	0	11	
1495	1472	0	3.0	0	4	1	20	
1618	502	0	0.8	0	7	0	52	
1099	1697	0	0.5	0	0	1	60	
1307	831	0	1.7	1	7	1	26	
...	
14	1866	0	0.5	0	13	1	52	
282	1839	1	1.2	0	9	1	54	
952	1444	1	2.1	0	9	0	38	
1079	1893	1	0.5	1	1	0	23	
486	1089	1	0.9	1	12	1	2	

	m_dep	mobile_wt	n_cores	pc	px_height	px_width	ram	sc_h	sc_w	\
423	0.4	158	2	13	195	1205	1122	12	6	
1495	0.3	169	2	6	443	892	797	6	1	
1618	1.0	82	6	8	281	1159	2666	5	4	
1099	0.1	90	4	0	88	1046	441	15	1	
1307	0.7	177	5	11	511	621	1704	6	5	
...	
14	0.7	185	1	17	356	563	373	14	9	
282	0.5	200	7	11	475	1493	927	19	10	
952	0.4	104	7	16	624	917	3764	14	9	
1079	0.1	179	8	3	1203	1432	1482	15	7	
486	0.7	145	5	15	636	1259	2765	13	12	
...										
1079	17	0		1	0					
486	10	1		0	1					

[600 rows x 20 columns]

```
print(x_train)
```

	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_memory	\
993	686	1	0.5	0	11	0	3	
1156	1732	0	0.8	0	2	0	61	
615	880	0	0.5	1	1	0	44	
703	1413	0	0.5	1	4	1	39	
1130	1975	1	1.9	1	2	0	31	
...	
1016	591	1	2.8	0	0	1	54	
165	517	0	1.4	1	3	1	33	
7	1954	0	0.5	1	0	0	24	
219	1551	0	1.1	0	4	0	51	
1350	1398	0	1.6	1	8	1	26	

	m_dep	mobile_wt	n_cores	pc	px_height	px_width	ram	sc_h	sc_w	\
993	0.3	91	6	15	1109	1392	570	7	6	
1156	0.3	172	5	3	201	656	3940	17	11	
615	0.5	172	8	15	436	1302	3132	8	7	
703	0.1	185	5	12	1039	1318	3878	19	16	
1130	0.9	151	1	17	775	1607	3022	13	5	
...	
1016	0.1	172	7	15	169	1916	1414	6	1	
165	0.8	183	4	8	660	974	3704	17	16	
7	0.8	187	4	0	512	1149	700	16	3	
219	0.1	88	5	6	1738	1995	3844	11	8	
1350	0.8	150	1	12	755	1284	3488	14	3	
...										
219	4	0		1	0					
1350	11	1		1	0					

```
from sklearn.preprocessing import MinMaxScaler
```

```
d=MinMaxScaler()
```

```
y=train['price_range']
x=train.drop('price_range',axis=1)
print(x)
print(y)
```

```

battery_power  blue  clock_speed  dual_sim  fc  four_g  int_memory  \
0           842    0           2.2         0  1         0           7
1          1021    1           0.5         1  0         1          53
2           563    1           0.5         1  2         1          41
3           615    1           2.5         0  0         0          10
4          1021    1           1.2         0  1         1          44
...          ...    ...         ...     ...  ..     ...         ...
1995         794    1           0.5         1  0         1           2
1996        1965    1           2.6         1  0         0          39
1997        1011    0           0.9         1  1         1          36
1998        1512    0           0.9         0  4         1          46
1999         510    1           2.0         1  5         1          45

m_dep  mobile_wt  n_cores  pc  px_height  px_width  ram  sc_h  sc_w  \
0      0.6       100       2  2         20       756  2549   9    7
1      0.7       136       3  6         905     1988  2631  17   3
2      0.9       145       5  0        1263     1716  2003  11   2
3      0.8       131       8  0        1216     1786  2769  16   8
19  wifi         2000  non-null  int64
20  price_range  2000  non-null  int64
dtypes: float64(2), int64(19)
memory usage: 328.2 KB

```

```
print(train.isnull().sum())
```

```

battery_power    0
blue             0
clock_speed      0
dual_sim         0
fc               0
four_g           0
int_memory       0
m_dep            0
mobile_wt        0
n_cores          0
pc               0
px_height        0
px_width         0
ram              0
sc_h             0
sc_w            0
talk_time        0
three_g          0
touch_screen     0
wifi             0
price_range      0
dtype: int64

```

std	0.499662	18.145715	0.288416	35.399655	2.287837	...
min	0.000000	2.000000	0.100000	80.000000	1.000000	...
25%	0.000000	16.000000	0.200000	109.000000	3.000000	...
50%	1.000000	32.000000	0.500000	141.000000	4.000000	...
75%	1.000000	48.000000	0.800000	170.000000	7.000000	...
max	1.000000	64.000000	1.000000	200.000000	8.000000	...

	px_height	px_width	ram	sc_h	sc_w	\
count	2000.000000	2000.000000	2000.000000	2000.000000	2000.000000	
mean	645.108000	1251.515500	2124.213000	12.306500	5.767000	
std	443.780811	432.199447	1084.732044	4.213245	4.356398	
min	0.000000	500.000000	256.000000	5.000000	0.000000	
...						
75%	16.000000	1.000000	1.000000	1.000000	2.250000	
max	20.000000	1.000000	1.000000	1.000000	3.000000	

[8 rows x 21 columns]

Output is truncated. View as a [scrollable element](#) or open in a [text editor](#). Adjust cell output [settings](#).

```
train.info()
```

[5]

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2000 entries, 0 to 1999
Data columns (total 21 columns):
#   Column      Non-Null Count  Dtype
---  ---
0   battery_power  2000 non-null   int64
1   blue          2000 non-null   int64
2   clock_speed    2000 non-null   float64
3   dual_sim       2000 non-null   int64
4   fc            2000 non-null   int64
..  ..
..  ..
```

```
import pandas as pd
```

```
train=pd.read_csv('/content/train.csv')
print(train.describe())
```

	battery_power	blue	clock_speed	dual_sim	fc	\
count	2000.000000	2000.0000	2000.000000	2000.000000	2000.000000	
mean	1238.518500	0.4950	1.522250	0.509500	4.309500	
std	439.418206	0.5001	0.816004	0.500035	4.341444	
min	501.000000	0.0000	0.500000	0.000000	0.000000	
25%	851.750000	0.0000	0.700000	0.000000	1.000000	
50%	1226.000000	0.0000	1.500000	1.000000	3.000000	
75%	1615.250000	1.0000	2.200000	1.000000	7.000000	
max	1998.000000	1.0000	3.000000	1.000000	19.000000	

	four_g	int_memory	m_dep	mobile_wt	n_cores	...	\
count	2000.000000	2000.000000	2000.000000	2000.000000	2000.000000	...	
mean	0.521500	32.046500	0.501750	140.249000	4.520500	...	
std	0.499662	18.145715	0.288416	35.399655	2.287837	...	
min	0.000000	2.000000	0.100000	80.000000	1.000000	...	
25%	0.000000	16.000000	0.200000	109.000000	3.000000	...	
50%	1.000000	32.000000	0.500000	141.000000	4.000000	...	
75%	1.000000	48.000000	0.800000	170.000000	7.000000	...	
max	1.000000	64.000000	1.000000	200.000000	8.000000	...	

	px_height	px_width	ram	sc_h	sc_w	\
count	2000.000000	2000.000000	2000.000000	2000.000000	2000.000000	
mean	645.108000	1251.515500	2124.213000	12.306500	5.767000	
std	443.780811	432.199447	1084.732044	4.213245	4.356398	

```
confusion=confusion_matrix(y_test,yp)
print(confusion)
```

[49]

```
... [[159   4   0   0]
      [  6 140   4   0]
      [  0   3 128   8]
      [  0   0   5 143]]
```